

Xueqi Cheng

xc25@fsu.edu · xueqic.github.io · github.com/xueqic · linkedin.com/in/xqc · Google Scholar

BIOGRAPHY

Xueqi Cheng is a Ph.D. student in Computer Science at Florida State University, advised by Dr. Yushun Dong, and a member of the Responsible AI (RAI) Lab. His current research interests include agentic AI (e.g., agent distillation and model routing), LLM compression (e.g., knowledge distillation, pruning, and quantization), and multimodal models. He is also interested in AI for social good and real-world applications of AI, including social network analysis and civil and infrastructure engineering.

EDUCATION

Florida State University

Doctor of Philosophy (**Ph.D.**) in Computer Science
• Advisor: Dr. Yushun Dong, Responsible AI (RAI) Lab

2025 – present

Vanderbilt University

Doctor of Philosophy (**Ph.D.**) in Computer Science (transferred out)
• Advisor: Dr. Tyler Derr, Network and Data Science (NDS) Lab

2023 – 2025

Master of Science (**M.S.**) in Systems Engineering

2021 – 2023

RESEARCH EXPERIENCE

Applied Research, Nokia of America Corporation

Research Intern

May 2026 – present

- Research topic: Agent distillation
- Manager: Kelly Wan; Mentor: Dr. Liang Wu

Responsible AI Lab, Florida State University

Ph.D. Student

Aug 2025 – present

- Research areas: Agentic AI, LLM compression, multimodal models
- Advisor: Dr. Yushun Dong

Chief Data Office, AT&T Labs

Research Intern

Jun 2025 – Aug 2025

- Research topic: Efficient LLM serving for multi-turn dialogues
- Manager: Dr. Divesh Srivastava; Mentors: Dr. Qiong Wu, Dr. Zhengyi Zhou

Network and Data Science Lab, Vanderbilt University

Ph.D. Student

Aug 2023 – May 2025

- Research areas: Geometric deep learning, social network analysis, AI for social good
- Advisor: Dr. Tyler Derr

SCHOLARSHIPS & AWARDS

- Osher Lifelong Learning Institute Scholarship, Florida State University Apr 2026
- WSDM NSF Travel Award Feb 2026
- CS Student Seminar Best Presentation Runner-Up, Florida State University Nov 2025
- Naaman Franklin Faile Jr. Graduate Fellowship, Florida State University Sep 2025
- Engineering Graduate Fellowship, Vanderbilt University Aug 2023
- IBM PhD Fellowship, Vanderbilt University Aug 2021

PUBLICATIONS

* denotes equal contribution (co-first authors).

Conference Papers

- [C04] **Xueqi Cheng**, Yushun Dong. LatentRouter: Can We Choose the Right Multimodal Large Language Model Before Seeing Its Answer? In Proceedings of the 40th Conference on Neural Information Processing Systems (NeurIPS), 2026. (acceptance rate 25.7%)

- [C03] **Xueqi Cheng**, Catherine Yang, Yuying Zhao, Yu Wang, Hamid Karimi, Tyler Derr. BTS: A Comprehensive Benchmark for Tie Strength Prediction. In Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD), 2025. (**Oral**; acceptance rate 22%)
- [C02] **Xueqi Cheng***, Yu Wang*, Yunchao (Lance) Liu, Yuying Zhao, Charu C. Aggarwal, Tyler Derr. Edge Classification on Graphs: New Directions in Topological Imbalance. In Proceedings of the 18th ACM International Conference on Web Search and Data Mining (WSDM), 2025. (acceptance rate 17.3%)
- [C01] Yu Wang, Tong Zhao, Yuying Zhao, Yunchao Liu, **Xueqi Cheng**, Neil Shah, Tyler Derr. A Topological Perspective on Demystifying GNN-based Link Prediction Performance. In Proceedings of the 12th International Conference on Learning Representations (ICLR), 2024. (acceptance rate 31%)

Journal Papers

- [J03] Bo Ni, Zheyuan Liu, Leyao Wang, Yongjia Lei, Yuying Zhao, **Xueqi Cheng**, Qingkai Zeng, Luna Dong, Yinglong Xia, Krishnaram Kenthapadi, Ryan Rossi, Franck Dernoncourt, Md Mehrab Tanjim, Nesreen Ahmed, Xiaorui Liu, Wenqi Fan, Erik Blasch, Yu Wang, Meng Jiang, Tyler Derr. Towards Trustworthy Retrieval Augmented Generation for Large Language Models: A Survey. ACM Computing Surveys (CSUR), 58(15), 2026.
- [J02] Yi Zhang, Yuying Zhao, Zhaoqing Li, **Xueqi Cheng**, Yu Wang, Olivera Kotevska, Philip S. Yu, Tyler Derr. A Survey on Privacy in Graph Neural Networks: Attacks, Preservation, and Applications. IEEE Transactions on Knowledge and Data Engineering (TKDE), 36(12), 2024. (impact factor 8.9)
- [J01] Yuying Zhao, Yu Wang, Yunchao Liu, **Xueqi Cheng**, Charu Aggarwal, Tyler Derr. Fairness and Diversity in Recommender Systems: A Survey. ACM Transactions on Intelligent Systems and Technology (TIST), 16(1), 2025. (impact factor 7.2)

Workshop and Symposium Papers

- [W02] **Xueqi Cheng**. Edge-centric Network Analytics. In Proceedings of the 18th ACM International Conference on Web Search and Data Mining (WSDM) Doctoral Consortium, 2025.
- [W01] **Xueqi Cheng**, Yan Zhang, Sining Lu, Yinqiao Zhu, Wei He. Urban Bridge Conceptual Design Based on Virtual Reality Graphic Engine. In IABSE Symposium Report, International Association for Bridge and Structural Engineering (IABSE), 2017.

Preprints

- [P06] **Xueqi Cheng***, Qinwen Ge*, Hamid Karimi, Yushun Dong, Tyler Derr. Adverse Online Social Interactions: A Multi-Level Evolutionary Analysis of Local Patterns, Diffusion, and Community Disruption. arXiv preprint arXiv:2606.20846, 2026.
- [P05] Runyang Xu*, **Xueqi Cheng***, Yushun Dong. A Nationwide Benchmark for Wildfire Initial Attack Failure Prediction with Public Environmental Data. arXiv preprint arXiv:2606.15529, 2026.
- [P04] **Xueqi Cheng**, Qiong Wu, Zhengyi Zhou, Xugui Zhou, Tyler Derr, Yushun Dong. SOMA: Efficient Multi-turn LLM Serving via Small Language Model. arXiv preprint arXiv:2605.11317, 2026.
- [P03] **Xueqi Cheng**, Xugui Zhou, Tyler Derr, Yushun Dong. ReAD: Reinforcement-Guided Capability Distillation for Large Language Models. arXiv preprint arXiv:2605.11290, 2026.
- [P02] **Xueqi Cheng**, Minxing Zheng, Shixiang Zhu, Yushun Dong. MISLEADER: Defending against Model Extraction with Ensembles of Distilled Models. arXiv preprint arXiv:2506.02362, 2025.
- [P01] Yuying Zhao, Yu Wang, **Xueqi Cheng**, Anne Marie Tumlin, Yunchao Liu, Damin Xia, Meng Jiang, Tyler Derr. Amplifying Your Social Media Presence: Personalized Influential Content Generation with LLMs. arXiv preprint arXiv:2505.01698, 2025.

OPEN-SOURCE PROJECTS

PyHazards

A Python framework for AI-powered hazard prediction and risk assessment. It provides a hazard-first API for loading data, building models, running end-to-end experiments, and extending with your own modules.

Jan 2026

TALKS	Industry Presentations	
	[I02] Efficient Context-dependent Multi-turn LLM Serving AT&T Labs	Aug 2025
	[I01] Predicting and Interpreting Churn Risks for Broadband Customers AT&T Labs	Jul 2025
	Guest Lectures	
	[L01] Introduction to Python/PyG CS 4352/5352: Social Network Analysis, Department of Computer Science, Vanderbilt University	Oct 2023
TEACHING EXPERIENCE	Vanderbilt University , Department of Computer Science	
	Teaching Assistant	Aug 2023 – May 2025
	• CS 4352/5352: Social Network Analysis • CS 4260: Introduction to Artificial Intelligence • CS 5892: Project in Data-Centric AI	
	Vanderbilt University , Department of Civil and Environmental Engineering	
	Teaching Assistant	Aug 2021 – May 2023
	• CE 2200: Statics • CE 3205: Structural Design • CE 3200: Structural Analysis	
MENTORING	Responsible AI Lab , Florida State University	
	Hugh Long, M.S. Computer Science	Fall 2025 – present
	Research topic: Wildfire prediction visualization	
	Runyang Xu, B.S. Statistics	Fall 2025 – present
	Research topic: Wildfire prediction foundation models; next position: M.S. Statistics, University of Michigan, Ann Arbor	
	Network and Data Science Lab , Vanderbilt University	
	Qinwen Ge, Ph.D. Computer Science	Spring 2025 – Spring 2026
	Research topic: Social network analysis	
	Bo Ni, Ph.D. Computer Science	Spring 2024 – Spring 2025
	Research topic: Trustworthy retrieval-augmented generation	
	Catherine Yang, M.S. Computer Science	Fall 2024 – Spring 2025
	Research topic: Social network analysis; next position: Software Engineer, Microsoft	
SERVICES	Organizing	
	• Publicity Chair, the 5th International Workshop on Machine Learning on Graphs (MLoG) at WSDM'24	
	Conference Reviewer / PC Member	
	• NeurIPS, ICML, KDD, ICWSM, WSDM, AAAI, WWW, SDM, IEEE Big Data, CIKM	
	Journal Reviewer	
	• Pattern Recognition, ACM Transactions on Knowledge Discovery from Data (TKDD)	